

Design and implementation of a programming language with LLVM

Josep Maria Domingo Catafal

Abstract— This article presents the design and development of a new programming language called *Craft*, using Rust and LLVM. The goal of the project is to gain a deeper understanding of how compilers work by creating one from scratch using these tools. The language is designed to be simple and easy to understand, but at the same time, it aims to be fast and efficient, so some sacrifices have to be made. The article covers the design and implementation of the language, including the lexer, parser, semantic analysis, and code generation. The project also includes a discussion of the challenges encountered during development and suggestions for future work. Overall, the project serves as a valuable learning experience for understanding the inner workings of compilers and the capabilities of Rust and LLVM.

Keywords— Programming Language, LLVM, SSA, Strongly Typed, Compiled

Resum— Aquest article presenta el disseny i desenvolupament d'un nou llenguatge de programació anomenat *Craft*, utilitzant Rust i LLVM. L'objectiu del projecte és aconseguir una millor comprensió de com funcionen els compiladors creant-ne un des de zero amb aquestes eines. El llenguatge està dissenyat per ser senzill i fàcil d'entendre, però al mateix temps, pretén ser ràpid i eficient, per la qual cosa s'han de fer alguns sacrificis. L'article cobreix el disseny i la implementació del llenguatge, incloent-hi el lexer, el parser, l'anàlisi semàntica i la generació de codi. El projecte també inclou una discussió sobre els reptes trobats durant el desenvolupament i suggeriments per a treballs futurs. En general, el projecte serveix com una valuosa experiència d'aprenentatge per comprendre el funcionament intern dels compiladors i les capacitats de Rust i LLVM.

Paraules clau— Llenguatge de programació, LLVM, SSA, Fortament tipat, Compilat

1 INTRODUCTION

HISTORICALLY, there has always been a dilemma between the speed of execution, and speed of development. Some languages are easy to program: they allow the programmer to not worry about low-level concepts such as memory management, and create abstractions that streamline the development. The problem is that these abstractions limit the language's efficiency, and create slower programs. Another reason that allows speeding up the development is dynamic typing, as it frees the user from the mental overhead that comes with deciding the type that should be used. But this comes with its own disadvantages, since it is very likely that you will encounter runtime errors unless you use something like a static analyzer, and even then errors can be frequent. These languages tend to be interpreted in order to save the programmer the time it takes to recompile the program, but it has an impact on the runtime performance of the program if we compare it to compiled languages. Python, for example, would be one of the largest representatives of this group of languages. On

the other hand we have languages like C, which have almost no abstraction and the programmer must be aware of what the computer is doing in every line of code he writes. They are languages that provide very good performance, but slow down the development, as the programmer must take into account many low-level concepts. An additional problem is that when managing memory manually, it opens the door to a lot of runtime errors in the form of memory leaks and segmentation faults. Currently, there are languages like Rust that solve these memory management issues without losing performance, but the development is still slow (maybe even slower) and the compilation time long. These languages tend to be strongly typed, which helps reduce runtime errors.

In this article we will explore the different types of programming languages, and we will implement a simple programming language that aims to find balance between the different trade-offs we talked about.

2 STATE OF THE ART

A programming language resembles, in a way, a natural language, and like natural languages, there's a lot of different ones with different properties. We mainly have two different ways to classify them: by paradigm and by how they

- Contact E-mail: jdomingocatafal@gmail.com
- Menció realitzada: Computació
- Tutorized by: Javier Sanchez Pujadas (Ciències de la Computació)
- Year 2022/23

are implemented. The paradigm allows us to classify them based on their features and the way we write programs with them. It tells us something about the design of the language. On the other hand we have the implementation which tells us information about how the programs are run. About how the code we write turns into instructions the CPU can execute. A programming language can have many implementations, but there's usually one that is the "official" one, and that's the ones we are going to focus on.

2.1 Paradigms

In the following sections we are going to take a look at the main paradigms in programming languages, even though, some of them may be considered sub-paradigms of the others. It's also important to note that most programming languages can fit into multiple of these paradigms, and thus they are multi-paradigm. We are going to discuss the most popular paradigms, but there are many more.

2.1.1 Imperative

Imperative programming languages follow a model of programming that is based on statements that change the state of the program. When we write an imperative program, we specify step by step what the computer has to do. Imperative programming is often implemented as procedural programming, a subtype in which statements are structured into procedures (also known as functions). Most of the programming languages we use nowadays can fit into this category, including C, C++, Java, C#, Go, Lua, JavaScript, etc.

2.1.2 Object-oriented

It's usually considered a subtype of imperative and procedural programming, even though there are languages in other paradigms that are also object-oriented. In object-oriented programming, data and behavior is grouped in units called objects. These objects hold data, and they can be modified by functions known as methods. The main representatives of this paradigm are languages like Java, C# or C++. These languages represent objects using classes and have many abstractions like inheritance and polymorphism.

2.1.3 Functional

Functional programming has the idea of treating computation as the evaluation of mathematical functions. Most of the functional languages are characterized by the use of recursive functions, immutable data and anonymous functions. Even though most modern languages incorporate functional programming in one way or another, some languages that we could call functional are Haskell (being one of the purest), OCaml, Clojure, Scala or Lisp, between others.

2.1.4 Logic

Logic programming is a programming paradigm that is based on formal logic. Some examples of logic programming languages include Prolog or Datalog.

2.2 Implementations

Programming languages can be implemented in different ways, but they generally fall into two main categories: compilers and interpreters. Both have it's upsides and it's downsides, and usually, depending on the design of the programming language and its goals, one may be a better option than the other.

2.2.1 Interpreted

Interpreted languages, are languages that don't need to be compiled before they are run. When you run a program they take the original source code and interpret it at runtime. The major downside of this type of implementations is that we may have made an error writing the code, and we won't find out until we run it. This can be quite dangerous since it means the program can crash at runtime because of a typo or some other kind of silly mistake that could've been caught by a compiler.

On the other hand, they offer a good developer experience, specially for prototyping or writing quick scripts. That's because you don't have to wait for the compiler to generate the code, you can just run it and make modifications to the program while it's running, without the need to restart it. They can also be embedded into other programs and be used to script the behavior of that program. An example of that are game engines, that are usually written in a compiled languages like C++, but they integrate a scripting language for all the gameplay development.

Some popular interpreted programming languages are Python, JavaScript, Lua or PHP. These are some of the most popular languages altogether, and that's because, they are very friendly languages, with a low barrier of entrance.

2.2.2 Compiled

Compiled languages are languages that turn the source code into machine code for a specific CPU architecture. They usually have a better runtime performance than interpreters, since they generate optimized machine code for that specific architecture. Another good thing about compiled languages is that they can catch a lot of bugs at compile time, since they have to check weather the program is valid in order to generate the machine code. This comes with a downside that is that it slows down development time, since we need to recompile the code on every change before we can run it.

Some popular compiled programming languages nowadays are C, C++, Go or Rust. There are different ways these compilers are implemented. Some of them, like Go, generate all the machine code by themselves. But others, like Rust or C/C++ with the clang compiler, generate an intermediate representation, and then use LLVM (we will look into LLVM in more detail later on) to generate the machine code. This way they don't have to re-implement the code generation, they just use the one LLVM does.

There's another breed of compiled languages, like for example Java or C#, that compile to bytecode, an intermediate instruction set, and then use a Virtual Machine that compiles Just In Time (JIT) the bytecode to machine code. These languages have the benefit, that they have interoperability with other languages that use that same Virtual Machine. For example in the case of Java, it uses the JVM

(Java Virtual Machine), and other languages like Kotlin or Clojure also use that same virtual machine, which makes them compatible with each other. Some other examples are Erlang with Elixir or C# with Visual Basic and F#.

3 GOALS

As we can see there's a lot of different paradigms with different goals and design choices that make them suitable for different things. The goal of this project is to design and implement a basic programming language that is user-friendly and that can be used to write simple programs. The language philosophy will revolve around the following ideas:

- **Compiled:** It will be a compiled language, in order to obtain a good performance and avoid run-time errors as much as possible. Compiled languages do a better job at catching programming errors.
- **Strongly typed:** Having a strong type system helps to avoid runtime errors, since a lot of the mistakes the programmer can make are caught at compile time. They also help document the code, making it easier to understand what it's doing. The downsides are that they may make the code more verbose and, sometimes, create a mental overhead when prototyping since the programmer has to think of the type to use. For bigger projects, though, this last downside is usually also present with dynamic typing, since dynamic typing will create a mental overhead when reading the code and trying to figure out the type of a variable.
- **Immutability:** Immutability will be enforced by default in order to avoid side effects, and to facilitate debugging. The user can still decide to make a variable mutable, but it has to be a conscious choice and not a given.
- **Easy to use:** The language has to be easy to use, giving a scripting-like feeling but with the benefits of a compiled strongly-typed language.

The new language doesn't aim to be an innovation, but rather a collection of ideas grabbed from other popular languages mixed together. The end goal is to learn about the internals of programming languages and better understand the design choices behind them. The language will be implemented using Rust and LLVM (more on that later), offering an opportunity to gain familiarity with LLVM, one of the most used technologies in the industry of compiler development, and to tackle a significant project from the ground up with Rust.

4 METHODOLOGY

Writing a compiler involves several steps. First, even though not strictly part of the compiler development, the language has to be designed. The syntax of the language, the symbols and the reserved keywords have to be defined. This design may evolve later on since we may find it's ambiguous or that it's too complex to implement.

Secondly we need to define the formal grammar for the language we designed. This will help us greatly when writing the parser. The grammar will tell us how the sentences of the language are built.

Thirdly, we have the implementation steps, which are three. The lexer, the parser and the semantic analysis and code generation. We have to implement them in that order, since they depend on the steps made by the previous one. We will see later on with more detail what these steps involve. But for now, simply we know that, for every new feature we want to add to the language, we have to implement it in every one of these steps.

It's also important to divide the languages into subsets, since this way we can start testing it. We don't need to implement all the features in the lexer, then in the parser, etc. We can implement the basic functionality in all the steps, and then start again on the lexer and add a new feature and so on. In the case of this project it was also important, since there's no time nor resources to implement all the features, so only a subset of features of the original design was implemented.

4.1 Git and GitHub

For version control Git was chosen since it's the industry standard and one of the most powerful tools out there. The repo is hosted on GitHub which is great for open source projects and allows us to use GitHub Actions, which run certain automated tasks we define on GitHub servers for free.

4.2 Quality Assurance

Apart from the implementation of the compiler itself, it's also important to test that it's actually working properly. To do that a test suite was developed, which consists of a few Craft programs, each testing different functionalities of the language, that are compiled and executed. Then the output of the program is checked and if it's not valid the test fails. This helps detect if any feature was broken during development, or if some edge cases were not taken into consideration.

Additionally, to leverage the features GitHub offers, a GitHub Action is run every time a Pull Request is opened. GitHub Actions, allow us to compile the project, run the test suite, run a linter and check that the formatting of the code follows the style guide. If any of these steps fails, the Pull Request cannot be merged to the main branch. This way we make sure that only working code (at least according to the test cases we defined) is merged into the stable branch.

The usual workflow when developing a new feature, would be to create a new branch, work on the code, and once done, a Pull Request to the master branch would be opened. Then the GitHub Actions would trigger, check everything is working as expected, and the code is formatted correctly. I would also review the code, and, if everything is correct, the Pull Request would be merged into master.

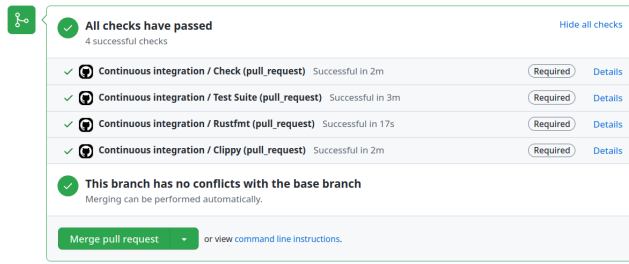


Fig. 1: Screenshot of a Pull Request where all checks finished successfully

5 LANGUAGE DESIGN

5.1 Primitive types

By default, Craft comes with some built-in (primitive), data types. These are the following:

- i64: 64 bit integer
- f64: 64 bit float
- string: array of characters
- bool: boolean

5.2 Start of the program

The starting point of craft programs is at the main function.

```
1 fn main() {
2   printf("Hello World!\n");
3 }
```

5.3 Functions

As you can see from the previous snippet, functions are declared with the `fn` keyword, followed by its name, and the parameters between parenthesis. The parameters are declared by specifying the identifier, followed by a colon and the type of the parameter. The return type comes after the closing parenthesis of the parameters. The braces that surround the body of the function are always mandatory.

```
1 fn something(a: i64, b: i64) i64 { ... }
```

5.3.1 Variable declaration

Variables are declared with the keyword `let`. The type of the variable is inferred from the value.

```
1 let four = 2 + 2; // type inferred to be an integer
2 let four = 2.0 + 2.0; // type inferred to be float
```

Variables are **immutable by default**, which means that once initialized their value cannot be changed. They act like constants that can be initialized at runtime. If the variable needs to be mutated, it can, but only by adding the keyword `mut` after `let`. So for example the following snippet would fail to compile:

```
1 fn main() {
2   let i = 0;
3   while i < 5 {
4     i = i + 1; // compile error, variable is not mutable!
5   }
6 }
```

The correct way to do it would be like this:

```
1 fn main() {
2   let mut i = 0;
3   while i < 5 {
4     i = i + 1; // compiles just fine, variable is mutable
5   }
6 }
```

This design choice comes from the idea that a lot of bugs come from mutating variables when we don't actually want to. By making them immutable by default, the user has to make the conscious choice of telling the compiler that he/she wants to mutate the variable. This way, if he/she mutates a variable he/she didn't intend to mutate, the compiler will catch it.

5.4 Loops and control flow

The way to repeat instructions multiple times in Craft is by using while loops. They can be used like so:

```
1 let i = 0;
2 while i < 100 {
3   // do stuff
4   i = i + 1;
5 }
```

Another way to repeat instructions is by using recursion:

```
1 fn fib(n: i64) i64 {
2   if n <= 1 {
3     n
4   } else {
5     fib(n - 1) + fib(n - 2)
6   }
7 }
```

Apart from loops, we also have if statements to create branches in our code. They can be used like so:

```
1 if a > b {
2   // do stuff
3 } else if a == b {
4   // do other stuff
5 } else {
6   // do other stuff
7 }
```

Notice that in both if and while statements, parenthesis around the conditional expression are not necessary. They can still be used if the programmer desires but the idiomatic way to do it is without parenthesis since it's easier to read and type.

5.5 Expressions and statements

Craft is an expression based language, which means that most constructs in the language are expressions. An expression is a valid unit of code that evaluates to a value. So for example, `2+2` is an expression, since it evaluates to a value (4 in this case). On the other hand we have statements, which don't evaluate to a value. For example, `(let a = 2+2;)` is a statement that contains the `2 + 2` expression, but the value is captured by the variable 'a', thus it's not returned, and it becomes a statement. Usually statements end in semicolon, except special cases like for example loops.

So when we said that Craft is expression based, it means that pretty much everything returns a value. For example if statements, are not actually statements, but expressions. We can do something like this:

```
1 let max = if a >= b { a } else { b };
```

In the code snippet above we are assigning the value returned from the `if` expression to a variable. We can do that, because the last item in each of the branches of the `if` expression is returning an expression (of the same type in both branches, else it's a compile time error). Whenever the last thing of a code block does not contain a semicolon, it means that block will return the value of evaluating that expression. So this actually works:

```
1 let two = {
2   2 + 2
3 };
4
5 let x = {
6   let x = func();
7   x
8 }
```

The block ends in an expression without semicolon, thus it returns the value of evaluating that expression (4) and assigns that to the variable. This works for functions too:

```
1 fn add(a: i64, b: i64) i64 {
2   a + b
3 }
4
5 fn max(a: i64, b: i64) i64 {
6   if a >= b { a } else { b }
7 }
```

If we were to add a semicolon to any of the return expressions on the previous functions, we would get a compilation error, since we would turn the expressions into statements, and the function would not return anything.

It's important to note that *if* expressions **must** have an else branch, since a value always has to be returned from the expression. If it's used as a statement (it doesn't return a value), then it's fine to have only the *if* branch.

Since blocks return values they can also be used for grouping expressions. For example, we can change the precedence of arithmetic operations using blocks:

```
1 let a = 2 * { 3 + 2 }; // a will equal 10
```

What the previous code does is evaluate 2, then evaluate the block which returns 5 and then multiply both expressions.

5.5.1 Examples of expressions

```
1 // function calls
2 func()
3 // arithmetic operations
4 2 + 2
5 // if else
6 if x != y { x } else { y }
7 // binary comparisons
8 x > y
9 // code blocks
10 { expr }
```

5.5.2 Examples of statements

```
1 let a = 2 + 2; // variable declaration
2 while true { } // loops
3 if a > b { printf("hello world"); } // ifs without return expr.
```

5.6 Return statements

We talked about blocks returning values by leaving the last expression without semicolon. So if we want a function to

return a value we just leave the last expression of the function without semicolon, and it will be returned. But what if you need to return earlier, maybe based on some condition. For those cases there's the `return` statement, which instantly returns from the function when called. It can be invoked with the **ret** keyword. For example, when implementing a recursive Fibonacci, we could do it like so:

```
1 fn fib(n: i64) i64 {
2   if n <= 1 {
3     ret n;
4   }
5   fib(n - 1) + fib(n - 2)
6 }
```

5.7 Data structures

We talked about some primitive data types, but that's not always enough to model more complex programs. That's why we can define custom data structures using structs. Structs work the same way they do in other languages like C or Rust. We declare them like so:

```
1 struct User {
2   name: string
3   age: i64
4 }
```

When creating an instance of a struct we do it the following way:

```
1 let user = User!{
2   name: "Tux",
3   age: 27,
4 };
```

5.8 Arrays

Craft has support for arrays. They are fixed size and they can be declared like so:

```
1 let a = [1, 2, 3] i64;
```

The type annotations for an array follow the pattern `[type; size]`. For example, when declaring a function that takes an int array as a parameter and returns that same array:

```
1 fn do_nothing(arr: [i64; 3]) [i64; 3] {
2   arr
3 }
```

Whenever we want to access an element of an array we use square brackets with the index of the element we want to access. Arrays in Craft are zero indexed, meaning the index of the first element is 0.

```
1 let a = [1, 2, 3] i64;
2 let b = 2 + a[0]; // b = 3
```

5.9 Print statements

In order to print to stdout with craft the `printf` statement can be used. It works the same way as `printf` in C (it actually calls `libc` under the hood).

```
1 let x = 2 + 2;
2 printf("This is 2 + 2 => %d\n", x);
```

5.10 Naming conventions

In Craft function and variable names are written in `snake_case`. It's only structs that are written in `PascalCase`.

6 TECH STACK

Writing a compiler does not require many tools, but some of them can help a lot in paving the road. For this compiler we are going to mainly use two tools: The Rust Programming Language and LLVM.

6.1 Rust

Rust is a systems programming language that was first released in 2010. It was developed by the Mozilla Foundation with the goal of creating a safe and concurrent language that would be suitable for low-level systems programming tasks, such as operating systems, and performance critical programs, like a browser engine. One of the major features of the language is it's guaranteed memory-safety (and thread-safety) without requiring the use of a garbage collector or reference counting (and thus not compromising on performance). This combined with its powerful type system, helps catch a lot of bugs at compile-time.

For this reason we are going to use this language. It's also really comfortable to use and comes with a lot of great tooling like cargo, which is the command line tool used for compiling, managing dependencies, etc. and clippy, a linter that comes built in and gives great hints on how to improve the code.

6.2 LLVM

LLVM is a toolchain for building compilers, i.e. a set of tools that help us implementing compilers. It was created in 2003 by Chris Lattner (also creator of the Swift programming language) and has the support of companies like Apple (LLVM is a part integral part of XCode and Swift for iOS application development), Google, IBM or Intel. Currently, there are several mainstream programming languages that use it, such as C and C++ (via the Clang compiler), Rust, Swift, Crystal...

As we said LLVM has a lot of tools, but among all them, the LLVM Core libraries are the most important and particularly relevant for us. We will be referring to them as LLVM from now on for simplicity.

LLVM will allow us to generate assembly for a lot of different architectures without any extra effort. It can even generate Web Assembly, which allows us to run the language in modern web browsers. The Craft compiler will generate LLVM IR (Intermediete Representation), and then it will be piped to LLVM, which will take the IR, apply transformations to it, in order to optimize it, and then the assembly for the target architecture will be generated.

Generating LLVM IR instead of assembly, also frees us from some headaches. For example, since LLVM is architecture independent, when we generate the code, we don't need to worry about the number of registers, as we have an unlimited number of virtual registers, which LLVM will later map to the registers of the corresponding architecture.

As we discussed, LLVM also applies optimizations to the generated code, like dead code elimination, constant folding, loop unrolling, etc. However, in order for LLVM to perform these optimizations, we will have to generate the code in SSA form (Static Single-Assignment). This means that we can only assign a value to a variable once. If we need to reassign a value, a new variable has to be created that replaces the other one. The reason SSA is used is that it makes applying optimizations a lot easier.

We will explore LLVM in more detail later in the article.

7 ARCHITECTURE

Most compilers are divided into two parts: the front-end and the back-end. The front-end is the part of the compilers that takes the source code and transforms it into an intermediate representation than will later be transformed into the actual machine code by the back-end. In our case, since we are using LLVM, we don't need to worry too much about the back-end, since LLVM will be in charge of generating the machine code. Our job will be to go from the source code to the LLVM intermediate representation (we will call it **IR** from now on). The front-end of the compiler is typically composed of four main steps: lexical analysis (implemented by the lexer), syntactic analysis (implemented by the parser), semantic analysis and the intermediate code generation.




Fig. 2: Diagram showing the compilation workflow of a Craft program

7.1 Lexer

This step consists of breaking the source code into a sequence of tokens. A token is a basic building block of the languages, such as a keyword or an identifier.

A lexer can be implemented using an external lexical analyzer, but in our case it will be implement by hand, since the goal is to learn how it works. Implementing a lexer from scratch is actually really simple. It works the following way:

1. Start by reading the source code character by character until we reach the end.
2. If the character by itself forms a valid sequence (e.g. a parenthesis) we create a token from it. If it doesn't, we continue reading characters until we find a valid sequence and then create the token.

Note that sometimes we may find a valid sequence, but that's not enough to create a token, since it may be the

start of another longer and valid sequence. We may need to check the following character/s to check whether it continues. An example of such case would be the '>' operator, since, by itself is a valid sequence, but it may be the start of the '>=' operator. So we need to check if the next character is an '=' or something else.

The lexer requires a bit of work to set up, but after that, expanding it is trivial, since we only need to add a new word to the list of reserved words, in the case we want to add a reserved word, or add a new rule that detects a new symbol for example.

7.2 Parser

The lexer allowed us to identify the symbols of the program, but it does not allow us to determine if their order is correct, or if they follow the rules of the language (i.e. it's grammatically correct). That is the job of the parser.

The parser takes the sequence of tokens obtained from the lexer and transforms it into an Abstract Syntax Tree (AST). This tree represents the structure of the program and determines its syntactic structure. It will tell us the order in which we need to execute the instructions.

To generate the AST, compilers use a context free grammar. A grammar is a set of rules that tells us how to form valid strings of tokens in a specific language. They are formed of a set of symbols, which can be divided into terminal and non-terminal symbols, and a set of production rules that specify how the non-terminal symbols can be replaced by sequences of terminal and non-terminal symbols. A context free grammar is a type of grammar which rules do not depend on the context in which the symbols appear.

The goal of the parser is to make the program obey the rules of the grammar. Here's an example of a simple grammar for parsing function prototypes:

```
<proto> ::= fn <id> "(" <params> ")"
<id> ::= letter {letter | digit | "_"}
<params> ::= "(" { <param> {, <param> } } ")"
<param> ::= <id>: <type>
```

In the example above, the things between the '<>' are the productions, and the things to the right of '::=' tell us how that rule is formed. For example the proto production tells us that the prototype of a function is formed by placing the fn keyword followed by an identifier (which is defined by the <id> production), an opening parenthesis, the params (defined by the <params> production) and a closing parenthesis. The curly braces indicate that the contents they surround can be repeated zero or more times.

If we were to translate the proto rule to code, it would look something like this:

```
1 fn parse_prototype() -> (Prototype, Err) {
2   // we expect to find the fn keyword,
3   // else it's an error
4   match current_token().kind {
5     // The advance function moves to the next token
6     TokenKind::Fn => advance(),
7     _ => return Err("Expected fn keyword"),
8   };
9
10  // we expect to find the function name,
11  // else it's an error
12  let name = match current_token().kind {
13    TokenKind::Identifier => current_token().lexeme,
14    _ => return Err("Expected an identifier"),
15  };
16 }
```

```
17 advance();
18
19 // we call the params rule
20 let params = parse_params();
21
22 // we are done, we return a struct
23 // with the info of the prototype
24 return Prototype {
25   name,
26   params,
27 };
28 }
```

7.3 Semantic Analysis and [IR] Code Generation

Once we have an AST, we can traverse it to generate the LLVM IR. Since this is a simple compiler, we are going to do the semantic analysis in this step. With more complex compilers, we may want to create a specific step of semantic analysis, but in our case it is not necessary.

LLVM has the concept of modules. A module contains all the information associated with one code file. If we have multiple files, we simply have to create different modules and link them. Modules contain functions and functions are made up of blocks. Blocks are defined by specifying a label. These labels are like assembly labels, they define sections within our code, and we can use them to make jumps to the different sections. Blocks are made up of instructions, similar to the instructions we find in assembly.




Fig. 3: Overview of the different building blocks of LLVM

Let's see an example of a small piece of code and what it would look like in the intermediate representation of LLVM. We have the following function that receives two integers and returns the maximum:

```
1 fn max(int a, int b) int {
2   if a > b { a } else { b }
3 }
```

If we translate it to LLVM IR we have the following code:

```
1 define i32 @max(i32 %a, i32 %b) {
2   entry:
3     %0 = icmp sgt i32 %a, %b
4     br i1 %0, label %btrue, label %bfalse
5
6   btrue:
7     br label %end
8
9   bfalse:
10    br label %end
11
12  end:
13    %retval = phi i32 [%a, %btrue], [%b, %bfalse]
14    ret i32 %retval
15 }
```


The first line of the previous code defines a function, which receives two 32 bits integers. A label entry is defined below.

The first thing it does, on the 'entry' tag, is comparing both integers (i.e. the if condition). The 'sgt' keyword, means 'signed greater than', which, as the names says, does a greater than signed comparison. The result of the comparison is a one bit integer which acts like a boolean. If it's true, it will jump to the label btrue, otherwise to the label bfalse.

In this case the two branches do the same thing: a jump to the label end. There we come across a concept called phi nodes. The phi nodes are kind of an inverted if. Depending on where we did the jump we will assign one value or another to the variable retval. If we come from btrue, retval is assigned %a, and if we come from bfalse it is assigned to %b.

Now, why do the two branches jump to the end tag and then do a conditional again in the end block? Couldn't we assign the value of retval directly inside the branch btrue or bfalse and spare us that third conditional? Well the answer is no, because then we would be generating code that is not in SSA form. And as we mentioned earlier, LLVM requires the generated code to be in SSA form (we would be assigning the value to %retval in two different places). And that's why phi nodes exist, to be able to solve this kind of problem.

7.3.1 LLVM API

It's important to understand the LLVM IR, but generating all of that code by hand is a lot of work and it would require us to put a lot of infrastructure in place. Fortunately, LLVM comes with a C++ API, that makes generating the IR a lot easier. But there's a little problem, we are not using C++. Thankfully, there's a Rust library called Inkwell[5], that offers a comfortable and idiomatic Rust interface to the LLVM API. Under the hood it calls the C++ API by using FFI (Foreign Function Interface).

The API is made of several entities. The most fundamental one is the **BasicBlock**, which represents the blocks that we find inside a function, i.e. the set of instructions we find within a label. If we group a bunch of BasicBlocks, we have an entity call **Function**, which, you guess it represents a function in the IR. And by grouping a bunch of functions we get the **Module** entity. As you may notice these entities represent the different parts that we saw when we talked about the IR.

Another important entity is **Context**. We just need an instance of it, and it will keep track of the state of the compilation.

Finally, we have the **Builder** entity, which is the one in charge of actually generating the code. For example, if we want to generate a comparison, like the one we saw on the IR example, we would call the **Builder**, telling it which instruction to generate:

```
Builder.CreateICmpSGT(a, b, "some-name")
```

The first two parameters are the two values being compared. The third indicates the name of the generated variable (we can give it any name we want, it's only useful for

debugging).

The previous call will generate the following code (which is the third line of the IR example we saw):

```
%nom = icmp sgt i32 %a, %b
```

8 RESULTS

The aim of this bachelor's thesis was to develop a programming language using Rust and LLVM. The language was designed to incorporate basic features such as if statements, while loops, functions, structs, and arrays.

The implementation of the language was carried out using Rust as the primary programming language and LLVM as the compiler infrastructure. The use of Rust allowed for a pleasant development experience and a fast compiler, while LLVM provided the necessary tools for code optimization and generation.

The final result of the project is a functioning programming language that includes the aforementioned features and can be used to write and execute basic programs. The implementation of if statements, while loops, functions, structs, and arrays was successful, and the language was able to perform as expected.

With the compiler as of today, we can write programs like the following:

```
1 struct ComplexNum {
2     real: f64
3     imaginary: f64
4 }
5
6 fn add(x: ComplexNum, y: ComplexNum) ComplexNum {
7     ComplexNum{
8         real: x.real + y.real,
9         imaginary: x.imaginary + y.imaginary,
10    }
11 }
12
13 fn main() {
14     let x = ComplexNum{
15         real: 1.0,
16         imaginary: 2.0,
17     };
18
19     let y = ComplexNum{
20         real: 3.0,
21         imaginary: 4.0,
22     };
23
24     let z = add(x, y);
25     printf("x + y = %f + %fi\n", z.real, z.imaginary);
26 }
```

Or this which prints an approximation of the Mandelbrot set:

```
1 fn mandelbrot(a: f64, b: f64) f64 {
2     let mut za = 0.0;
3     let mut zb = 0.0;
4
5     let mut i = 0;
6
7     while i < 50 {
8         za = (za*za - zb*zb) + a;
9         zb = (za*zb + za*zb) + b;
10        i = i+1;
11    }
12
13    za*za + zb*zb
14 }
15
16 fn main() {
17     let xstart = -2.0;
18     let xend = 0.5;
19     let ystart = 1.0;
20     let yend = -1.0;
21
22     let xstep = 0.0315;
23     let ystep = -0.05;
```



```

24
25 let mut x = xstart;
26 let mut y = ystart;
27
28 while y > yend {
29     x = xstart;
30     while x < xend {
31         if mandelbrot(x, y) < 4.0 {
32             printf("x");
33         } else {
34             printf(" ");
35         }
36         x = x + xstep;
37     }
38     printf("\n");
39     y = y + ystep;
40 }
41 }

```

The output of this program can be found on the appendix A.1.

Furthermore, the project provided an opportunity to gain experience in using Rust and LLVM, as well as to gain a deeper understanding of the design and implementation of programming languages. The combination of Rust and LLVM proved to be a powerful combination for developing a compiler.

9 CONCLUSIONS AND FUTURE WORK

The development of a programming language is a long and arduous process. It requires the work of a lot of very good engineers, and a lot of time. The goal of this project was not to create a production ready language packed with features and an extensive standard library, but rather to create a simple language with the basic functionality to make something useful. It was a project to learn about the internals of compilers and programming language design. And I think the goal was accomplished. We have a simple language that allows us to create programs that actually run. However, there are many opportunities for future work and improvement, since, as stated, it takes a long time to develop a good language:

- **New features:** The language presented in this paper, is very limited. It would be interesting to add new features such as for loops with ranges, a pipe operator similar to the one in Elixir, list comprehension like in Python, and many more features that would make the development process a lot more pleasant.
- **Modules:** As of now, all code has to go on one file, there's no way to import code from another one. A module system has to be designed and implemented in order to allow it.
- **Error handling mechanism:** Currently there's no way to handle errors. An error handling mechanism has to be put in place. The idea would be to use something similar to the error handling mechanism in Rust, rather than exceptions, since it gives the code linearity, and they have to be explicitly handled on each case.
- **Tooling:** Some extra tooling like a formatter or a linter could be included with the compiler.

REFERENCES

- [1] *LLVM Kaleidoscope Tutorial: Implementing a Language with LLVM*. Accessed the 25th of October 2022, from <https://llvm.org/docs/tutorial>
- [2] Rathi, M. *A complete guide to LLVM for programming language creators*. Mukuls Blogs. Accessed the 2nd of November 2022, from <https://mukulrathi.com/create-your-own-programming-language/llvm-ir-cpp-api-tutorial>
- [3] Nystrom, R. *Crafting Interpreters*. Published July 2021. ISBN 0990582930.
- [4] *Mapping High Level Constructs to LLVM IR*. Accessed the 20th of October 2022 from <https://mapping-high-level-constructs-to-llvm-ir.readthedocs.io/en/latest/README.html>
- [5] *Inkwell Documentation*. Accessed the 4th of October 2022 from <https://thedan64.github.io/inkwell/inkwell/index.html>
- [6] Ball, T. *Writing A Compiler In Go*. Published August 2018. ISBN 398201610X.

A.1 Output of the Mandelbrot set program

[illegible]